

Beyond Cluely: 24/7 Organizational Intelligence

A 20-phase, 153-day governed execution policy integrating the Live Copilot and Ambient Intelligence roadmaps. Every phase verified by execution. The anti-Cluely: ambient, deep, rich, learning, evidence-backed.

Document	Maestro Ambient Intelligence — Governed Execution Policy
Audience	CEO, Engineering, Auditor, Investors
Date	2026-07-07
Current HEAD	d9026e2 (origin/main)
Total scope	20 phases, 153 days, ~612 hours, \$480K-720K
Source	MAESTRO_AMBIENT_INTELLIGENCE_ROADMAP.md (1111 lines, 12 phases)
Companion	MAESTRO_LIVE_COPILOT_ROADMAP_GOVERNED.md (8 phases)
CEO promise	Ambient org intelligence layer — works 24/7, not just during calls

GOVERNANCE LOOP READ RECEIPT (Coder) — 2026-07-07T05:50:00Z
GOVERNANCE_LOOP.md: "THE LOOP CANNOT BE BROKEN." · P20: "If M of N call sites pass it, the fix is (M/N)% done." · P26: "Principles don't enforce themselves. Re-application does." · P27: "Read at least the key assertions." · P34: "The auditor's method is not a memory — it's a checklist."
> Auditor: Read GOVERNANCE_LOOP.md, AUDITOR_GOVERNANCE.md, and ENTROPY_RECOVERY.md Part Four + Part Five from disk before auditing.

Table of Contents

Table of Contents	2
Chapter 1 — The Non-Negotiable Ethical Line	3
Chapter 1 — Architecture Overview	4
Chapter 2 — The CEO's Framing: Deeper, Richer, Ambient	5
2.1 DEEPER — Multi-Layer Intelligence	5
2.2 RICHER — Full Context, Not Just Audio	5
2.3 AMBIENT — Always Working, Not Just During Calls	5
2.4 The 5 Differences From Cluely (the moat)	5
2.5 The Moat (Why This Cannot Be Cloned)	6
2.6 The Bottom Line (the category claim)	6
Chapter 3 — The 20-Phase Unified Plan (phases map to the 3 dimensions)	7
3.1 Live Copilot Phases (Days 1-33)	7
3.2 Ambient Intelligence Phases (Days 34-153)	7
Chapter 4 — The Coding Execution Policy	8
4.1 Pre-Phase (BLOCKING — before any code)	8
4.2 During-Phase (per commit)	8
4.3 Post-Phase (BLOCKING — before next phase)	8
4.4 Forbidden Anti-Patterns (auditor rejects on sight)	8
Chapter 5 — Ambient Phase Details (9-20)	10
Chapter 6 — The Loop + Honest Disclosure	11
6.1 Honest Disclosure	11
6.2 Read Receipt	11

Chapter 1 — The Non-Negotiable Ethical Line

The Ambient Intelligence layer extends the Live Copilot ethical line. The differentiator from Cluely is absolute: Maestro is ambient (works between calls), deep (multi-layer), rich (full context), learning (compounds), evidence-backed (every suggestion cites organizational data). Cluely helps you cheat in the moment. Maestro helps your organization learn.

The following code paths are FORBIDDEN. Any appearance rejects the entire phase:

Forbidden path	Why rejected
Hiding from meeting software	Stealth framing violates the bright line
Bypassing screen-share detection	Deception; enterprise buyers will not tolerate it
Unconsented recording (audio/video/email/Slack/calendar)	Consent MUST precede every capture
Exam/interview assistance	Manipulation; rejected absolutely
"Undetectable mode" / stealth framing	Side panel is transparent, not hidden
Auto answer injection into calls	Suggestions display; the user speaks them
Inferring commitments from ambiguous language	Only explicit commitments tracked
Reading email/Slack/calendar content without channel consent	Channel consent, revocable
Emotion analysis used to manipulate	For user awareness only, never to "win"

Chapter 1 — Architecture Overview

The Ambient Intelligence Layer sits above the Organizational Memory (SituationSnapshot 27 fields + OutcomeLedger + OEM signals) and below the user-facing surfaces (Today panel + Whisper push + Ask + cross-meeting threads). It fuses signals from Calendar, Email, Slack/Teams, CRM connectors, and the Live Copilot audio stream into a 24/7 background loop: ingest → classify → correlate → alert.

The fusion layer feeds four intelligence engines: Commitment Escalation, Deal Health, Relationship Dynamics, and Sentiment & Emotion. Each engine writes back to the Organizational Memory, so every interaction makes the system smarter. This is the moat: Cluely has GPT; Maestro has your organization's entire history, learning from every interaction, building institutional memory that compounds over time.

The design law (from the constitution): every increase in internal intelligence must reduce external complexity. The ambient layer makes the user's day simpler (fewer forgotten commitments, fewer unprepared meetings, fewer repeated mistakes) even as the intelligence gets deeper. The UI stays at 4 sidebar items; the engines multiply invisibly behind them.

Chapter 2 — The CEO's Framing: Deeper, Richer, Ambient

The CEO's promise to investors is not "Cluely with better evidence." It is a fundamentally different category. Every phase in this plan must be evaluable against the 3 dimensions and the 5 differences. If a phase does not advance at least one dimension, it is theater and the auditor rejects it.

2.1 DEEPER — Multi-Layer Intelligence

Layer	What it does	Example	Phase
Sentiment Tracking	Real-time emotion from voice tone	Sam's voice stress +40% on pricing	10
Negotiation Strategy	BATNA, anchoring detection	Anchored \$50K. BATNA \$65K. Counter \$58K.	12
Relationship Dynamics	Power mapping, influence	Raj works to Sam on pricing. Address Sam.	13
Deal Health Scoring	Live probability during call	Deal health 78% to 62% after objection	11
Cross-Meeting Threads	Conversation continuity	4th time pricing came up. 2 closed, 1 lost.	14

2.2 RICHER — Full Context, Not Just Audio

Source	What it adds	Phase
Calendar	3 Globex meetings this week — deal acceleration	9
Email	Sam's last email 5 days ago, no response — risk	17
Slack	Raj asked about circuit breaker — mention it	17
CRM	Deal stage: Negotiation. \$3.2M at risk. Renewal 47 days	9, 11
Commitments	Promised pricing by Thursday. It's Friday. Overdue.	9

2.3 AMBIENT — Always Working, Not Just During Calls

When	What Maestro does	Phase
Between meetings	Globex call in 2 hours. Pre-call brief, 3 talking points.	9
After work hours	3 commitments overdue. Priority list with drafts.	9, 19
Weekend	Initech critical. Last interaction 22 days. Re-engage draft.	13, 19
Monday morning	8 meetings, 3 at-risk, 5 overdue. Priority matrix.	9, 20

2.4 The 5 Differences From Cluely (the moat)

#	Difference	Phases
1	Works between calls (24/7, not just during)	9, 19

#	Difference	Phases
2	Learns from every interaction (Call 1 = 40%, Call 10 = 89%)	5, 16, 20
3	Predicts failures before they happen (73% failure probability)	9
4	Connects everything (calendar + email + Slack + CRM + commitments)	9, 17
5	Builds institutional memory (calls → signals → LOs → patterns → laws)	All

2.5 The Moat (Why This Cannot Be Cloned)

Cluely's moat: GPT-4 API + audio capture. Anyone can clone this in a weekend.

Maestro's moat: Your organization's entire execution history, learning from every interaction, building institutional memory that compounds over time.

Cluely cannot say "You promised SSO by Day 60 — today is Day 45" (no commitment tracking). Cluely cannot say "The last 3 times we faced this pricing objection with a healthcare buyer, the deal closed when we offered a pilot" (no historical data). Cluely cannot say "Priya is on PTO starting Wednesday — delegate approvals today" (no team calendar). Cluely cannot say "Relationship with Initech is critical — last interaction 22 days ago" (no email/Slack monitoring). Maestro can. Because Maestro has the data. Cluely just has GPT.

2.6 The Bottom Line (the category claim)

This roadmap does not add "Cluely-like features" to Maestro. It creates a new category: ambient organizational intelligence.

Cluely = teleprompter for individuals.

Maestro Ambient = institutional memory for organizations.

One helps you cheat in the moment. The other helps your organization learn, predict, and compound intelligence over time. Every phase must advance this category claim. If a phase is "Cluely-like" (single-user, in-the-moment, no learning), the auditor rejects it.

Chapter 3 — The 20-Phase Unified Plan (phases map to the 3 dimensions)

The Ambient Intelligence roadmap (12 phases) integrates with the Live Copilot roadmap (8 phases). The Live Copilot phases (1-8) deliver the meeting-time layer. The Ambient phases (9-20) deliver the always-on layer. Together: 20 phases, 153 days, ~612 hours.

3.1 Live Copilot Phases (Days 1-33)

Phase	Days	Hours	Deliverable
1: Extension scaffold	1-3	12	manifest.json, consent-manager, panel shell
2: Audio + transcription	4-7	16	offscreen audio capture, Whisper STT
3: Scene 1 pre-call	8-11	16	lobby detection, attendee intelligence
4: Scene 2 live	12-18	28	4 card types (objection/commitment/whisper/pattern)
5: Scene 3 post-call	19-23	20	summary, draft email, "What Maestro learned"
6: Evidence + confidence	24-27	16	evidence-chain links, P25 confidence gate
7: Accessibility + polish	28-30	12	keyboard nav, aria-live, contrast
8: Integration + audit	31-33	12	E2E test, cross-surface coherence

3.2 Ambient Intelligence Phases (Days 34-153)

Phase	Days	Hrs	Deliverable	Key File
9: Ambient signal fusion	34-43	40	Calendar awareness + commitment escalation	calendar_awareness.py, commitment_escalation.py
10: Sentiment & emotion	44-53	40	Voice tone analysis, emotion detection	sentiment_engine.py
11: Deal health score	54-63	40	Live scoring, risk factors, momentum	deal_health.py
12: Negotiation strategy	64-73	40	BATNA, anchoring, concessions	negotiation_strategy.py
13: Relationship dynamics	74-83	40	Influence networks, power, coalitions	relationship_dynamics.py
14: Cross-meeting threads	84-93	40	Continuity, topic evolution, decisions	cross_meeting_threads.py
15: Talk ratio + coach	94-103	40	Speaking time, interruptions, clarity	talk_ratio_coach.py
16: Meeting grade + analysis	104-113	40	Effectiveness score, action items	meeting_grader.py
17: Email/Slack signals	114-123	40	Written comms monitoring, response time	written_signal_fusion.py
18: Multi-language	124-133	40	Accent-aware STT, cultural context, translation	translation_support.py
19: Ambient notifications	134-143	40	Smart nudges, context timing, DND	ambient_notifications.py
20: Advanced analytics	144-153	40	Trends, team performance, org learning	advanced_analytics.py

Totals: 20 phases, 153 days, ~612 hours, \$480K-720K (4 engineers × 6 months).

Chapter 4 — The Coding Execution Policy

This is the governed coding policy. Every phase **MUST** follow it. The auditor verifies compliance by execution (P31). The policy has 3 stages: pre-phase (blocking), during-phase (per commit), and post-phase (blocking).

4.1 Pre-Phase (BLOCKING — before any code)

Step	Action	Principle
1	Paste complete 8-field read receipt	Gate (CEO rejects without it)
2	Verify prior phase gate passes by execution	P1, P31
3	List files + grep call sites for parameter changes	P20
4	Audit every capture path for consent gating	Ethical line

4.2 During-Phase (per commit)

Step	Action	Principle
1	Commit includes VERIFICATION section with pasted output	P23
2	grep call sites; M of N = (M/N)% done	P20, Gate 15
3	Write 2 tests: unit + integration (production path)	P22
4	Re-run SSO scenario after shared-component changes	P24, P29
5	Gate every confidence value on denominator (< 10 = "insufficient")	P25
6	Execute consent-denied path; confirm zero capture	Ethical line

4.3 Post-Phase (BLOCKING — before next phase)

Step	Action	Principle
1	Run phase gate commands; paste output	P1
2	Run L0 gate; verify no regression	P29, P30
3	Push to origin/main; paste HEAD + origin/main (must match)	Gate 11
4	Auditor independently verifies by execution	P31
5	Only "Phase N PASS" allows next phase	The loop

4.4 Forbidden Anti-Patterns (auditor rejects on sight)

Anti-pattern	Principle violated
Theater tests (assert True, weak isinstance)	P27

Anti-pattern	Principle violated
Self-certification (✓ VERIFIED without output)	P5, P23
Wiring gaps (0 of N callers pass new param)	P20, Gate 15
Stale-clone auditing (HEAD != origin/main)	Gate 11
Single-input testing (only the golden case)	P28
Confidence without denominator	P25
Stealth framing (hide from meeting software)	Ethical line

Chapter 5 — Ambient Phase Details (9-20)

Each ambient phase has deliverables, an ethical guard, and a gate. The gates are verified by execution (P31). The ethical guards extend the Live Copilot bright line to the always-on layer: consent for every capture, no manipulation, no surveillance, no inference without confirmation.

Phase	Deliverable	Ethical Guard	Gate
9	Calendar awareness + commitment escalation	Calendar data only; no content with test cases; explicit consent only	test_calendar_explicit_test_only; commitment_escalation.py
10	Sentiment & emotion tracking	For user awareness only; never shown to other party	test_sentiment_engine.py (emotion never shown to other party)
11	Deal health score	Score with denominator (P25); < 10 deals insufficiently (confidence has denominator)	test_deal_health.py (confidence has denominator)
12	Negotiation strategy	Preparation only; no manipulative tactics; evidence chain required	test_negotiated_strategy.py (evidence chain required)
13	Relationship dynamics	Derived from org signals; no personal probes	test_relationship_dynamics.py
14	Cross-meeting threads	Threads by entity + topic; decisions tracked	test_cross_meeting_threads.py
15	Talk ratio + comms coach	Capability-building, not dominance-building	test_talk_ratio_coach.py (no dominance suggestions)
16	Meeting grade + analytics	Grade with denominator; follow-up tracking	test_meeting_grade.py
17	Email/Slack signal fusion	Per-channel consent, revocable; consent-denied path = 0	test_email_slack_fusion.py (consent-denied path = 0)
18	Multi-language support	Translation for user understanding; auto-translate support type (auto-translate gated)	test_multilingual_support.py (auto-translate gated)
19	Ambient notifications	Respects DND + focus + off-hours; max 5/day	test_ambient_notifications.py
20	Advanced analytics	Team-level aggregate; no individual surveillance	test_advanced_analytics.py (no individual PII)

Chapter 6 — The Loop + Honest Disclosure

The governance loop is the enforcement mechanism. It is the reason the 9 code-quality findings from Audit 1 are actually fixed (verified by the auditor at HEAD 09b2b87), not just claimed fixed. The loop applies equally to all 20 phases:

Step	Action	Who
1. Before each phase	Read governance modules from disk; paste 8-field receipt	Coder
2. During each phase	Gate P-numbers (P20, P22, P23, P24, P25) per commit	Coder
3. After each phase	Run gate + L0; push; paste HEAD + origin/main	Coder
4. Auditor verifies	Fetch → checkout → run gate independently (P31) → SSO scenario (P20)	Auditor
5. Next phase	Only after "Phase N PASS"	Both

6.1 Honest Disclosure

This is a 153-day, 612-hour, \$480K-720K build. The coder can execute every CODE phase (1-20). The CEO must resource the GTM milestones (design partners, ARR, SOC2/HIPAA/GDPR) in parallel. The 20 phases deliver the product the CEO promised investors: an ambient organizational intelligence layer that works 24/7, learns from every interaction, and compounds institutional memory over time. This is not a feature. This is a category.

6.2 Read Receipt

GOVERNANCE LOOP READ RECEIPT (Coder):

- GOVERNANCE_LOOP.md read at 2026-07-07T05:50:00Z. Key line: "THE LOOP CANNOT BE BROKEN."
- ENTROPY_RECOVERY.md Part Four (P20-P26) read. P20: "If M of N call sites pass it, the fix is (M/N)% done." P26: "Principles don't enforce themselves. Re-application does."
- ENTROPY_RECOVERY.md Part Five (P27-P34) read. P27: "Read at least the key assertions." P34: "The auditor's method is not a memory — it's a checklist."
- AUDITOR_GOVERNANCE.md Gates 15-20 read. Gate 15: "A function signature with a new parameter proves nothing." Gate 17: "Reading the dedup logic proves nothing; executing it does."
- audit_scripts/audit_gates.sh read — enforces Gate 11 + full suite + all verify scripts.

Auditor: Read GOVERNANCE_LOOP.md, AUDITOR_GOVERNANCE.md, and ENTROPY_RECOVERY.md Part Four + Part Five from disk before auditing. Do not audit from memory. The loop cannot be broken.