

# Beyond Cluely: 24/7 Organizational Intelligence

A 20-phase, 153-day governed execution policy integrating the Live Copilot and Ambient Intelligence roadmaps. Every phase verified by execution. The anti-Cluely: ambient, deep, rich, learning, evidence-backed.

|              |                                                                    |
|--------------|--------------------------------------------------------------------|
| Document     | Maestro Ambient Intelligence — Governed Execution Policy           |
| Audience     | CEO, Engineering, Auditor, Investors                               |
| Date         | 2026-07-07                                                         |
| Current HEAD | d9026e2 (origin/main)                                              |
| Total scope  | 20 phases, 153 days, ~612 hours, \$480K-720K                       |
| Source       | MAESTRO_AMBIENT_INTELLIGENCE_ROADMAP.md (1111 lines, 12 phases)    |
| Companion    | MAESTRO_LIVE_COPILOT_ROADMAP_GOVERNED.md (8 phases)                |
| CEO promise  | Ambient org intelligence layer — works 24/7, not just during calls |

GOVERNANCE LOOP READ RECEIPT (Coder) — 2026-07-07T05:50:00Z  
GOVERNANCE\_LOOP.md: "THE LOOP CANNOT BE BROKEN." · P20: "If M of N call sites pass it, the fix is (M/N)% done." · P26: "Principles don't enforce themselves. Re-application does." · P27: "Read at least the key assertions." · P34: "The auditor's method is not a memory — it's a checklist."  
> Auditor: Read GOVERNANCE\_LOOP.md, AUDITOR\_GOVERNANCE.md, and ENTROPY\_RECOVERY.md Part Four + Part Five from disk before auditing.

---

# Table of Contents

|                                                                  |          |
|------------------------------------------------------------------|----------|
| <b>Table of Contents</b>                                         | <b>2</b> |
| <b>Chapter 1 — The Non-Negotiable Ethical Line</b>               | <b>3</b> |
| <b>Chapter 2 — Architecture Overview</b>                         | <b>4</b> |
| <b>Chapter 3 — The 20-Phase Unified Plan</b>                     | <b>5</b> |
| 3.1 Live Copilot Phases (Days 1-33) . . . . .                    | 5        |
| 3.2 Ambient Intelligence Phases (Days 34-153) . . . . .          | 5        |
| <b>Chapter 4 — The Coding Execution Policy</b>                   | <b>6</b> |
| 4.1 Pre-Phase (BLOCKING — before any code) . . . . .             | 6        |
| 4.2 During-Phase (per commit) . . . . .                          | 6        |
| 4.3 Post-Phase (BLOCKING — before next phase) . . . . .          | 6        |
| 4.4 Forbidden Anti-Patterns (auditor rejects on sight) . . . . . | 6        |
| <b>Chapter 5 — Ambient Phase Details (9-20)</b>                  | <b>8</b> |
| <b>Chapter 6 — The Loop + Honest Disclosure</b>                  | <b>9</b> |
| 6.1 Honest Disclosure . . . . .                                  | 9        |
| 6.2 Read Receipt . . . . .                                       | 9        |

# Chapter 1 — The Non-Negotiable Ethical Line

The Ambient Intelligence layer extends the Live Copilot ethical line. The differentiator from Cluely is absolute: Maestro is ambient (works between calls), deep (multi-layer), rich (full context), learning (compounds), evidence-backed (every suggestion cites organizational data). Cluely helps you cheat in the moment. Maestro helps your organization learn.

The following code paths are FORBIDDEN. Any appearance rejects the entire phase:

| Forbidden path                                               | Why rejected                                      |
|--------------------------------------------------------------|---------------------------------------------------|
| Hiding from meeting software                                 | Stealth framing violates the bright line          |
| Bypassing screen-share detection                             | Deception; enterprise buyers will not tolerate it |
| Unconsented recording (audio/video/email/Slack/calendar)     | Consent MUST precede every capture                |
| Exam/interview assistance                                    | Manipulation; rejected absolutely                 |
| "Undetectable mode" / stealth framing                        | Side panel is transparent, not hidden             |
| Auto answer injection into calls                             | Suggestions display; the user speaks them         |
| Inferring commitments from ambiguous language                | Only explicit commitments tracked                 |
| Reading email/Slack/calendar content without channel consent | Channel consent, revocable                        |
| Emotion analysis used to manipulate                          | For user awareness only, never to "win"           |

---

## Chapter 2 — Architecture Overview

The Ambient Intelligence Layer sits above the Organizational Memory (SituationSnapshot 27 fields + OutcomeLedger + OEM signals) and below the user-facing surfaces (Today panel + Whisper push + Ask + cross-meeting threads). It fuses signals from Calendar, Email, Slack/Teams, CRM connectors, and the Live Copilot audio stream into a 24/7 background loop: ingest → classify → correlate → alert.

The fusion layer feeds four intelligence engines: Commitment Escalation, Deal Health, Relationship Dynamics, and Sentiment & Emotion. Each engine writes back to the Organizational Memory, so every interaction makes the system smarter. This is the moat: Cluely has GPT; Maestro has your organization's entire history, learning from every interaction, building institutional memory that compounds over time.

The design law (from the constitution): every increase in internal intelligence must reduce external complexity. The ambient layer makes the user's day simpler (fewer forgotten commitments, fewer unprepared meetings, fewer repeated mistakes) even as the intelligence gets deeper. The UI stays at 4 sidebar items; the engines multiply invisibly behind them.

# Chapter 3 — The 20-Phase Unified Plan

The Ambient Intelligence roadmap (12 phases) integrates with the Live Copilot roadmap (8 phases). The Live Copilot phases (1-8) deliver the meeting-time layer. The Ambient phases (9-20) deliver the always-on layer. Together: 20 phases, 153 days, ~612 hours.

## 3.1 Live Copilot Phases (Days 1-33)

| Phase                     | Days  | Hours | Deliverable                                         |
|---------------------------|-------|-------|-----------------------------------------------------|
| 1: Extension scaffold     | 1-3   | 12    | manifest.json, consent-manager, panel shell         |
| 2: Audio + transcription  | 4-7   | 16    | offscreen audio capture, Whisper STT                |
| 3: Scene 1 pre-call       | 8-11  | 16    | lobby detection, attendee intelligence              |
| 4: Scene 2 live           | 12-18 | 28    | 4 card types (objection/commitment/whisper/pattern) |
| 5: Scene 3 post-call      | 19-23 | 20    | summary, draft email, "What Maestro learned"        |
| 6: Evidence + confidence  | 24-27 | 16    | evidence-chain links, P25 confidence gate           |
| 7: Accessibility + polish | 28-30 | 12    | keyboard nav, aria-live, contrast                   |
| 8: Integration + audit    | 31-33 | 12    | E2E test, cross-surface coherence                   |

## 3.2 Ambient Intelligence Phases (Days 34-153)

| Phase                        | Days    | Hrs | Deliverable                                     | Key File                                        |
|------------------------------|---------|-----|-------------------------------------------------|-------------------------------------------------|
| 9: Ambient signal fusion     | 34-43   | 40  | Calendar awareness + commitment escalation      | calendar_awareness.py, commitment_escalation.py |
| 10: Sentiment & emotion      | 44-53   | 40  | Voice tone analysis, emotion detection          | sentiment_engine.py                             |
| 11: Deal health score        | 54-63   | 40  | Live scoring, risk factors, momentum            | deal_health.py                                  |
| 12: Negotiation strategy     | 64-73   | 40  | BATNA, anchoring, concessions                   | negotiation_strategy.py                         |
| 13: Relationship dynamics    | 74-83   | 40  | Influence networks, power, coalitions           | relationship_dynamics.py                        |
| 14: Cross-meeting threads    | 84-93   | 40  | Continuity, topic evolution, decisions          | cross_meeting_threads.py                        |
| 15: Talk ratio + coach       | 94-103  | 40  | Speaking time, interruptions, clarity           | talk_ratio_coach.py                             |
| 16: Meeting grade + analysis | 104-113 | 40  | Effectiveness score, action items               | meeting_grader.py                               |
| 17: Email/Slack signals      | 114-123 | 40  | Written comms monitoring, response time         | written_signal_fusion.py                        |
| 18: Multi-language           | 124-133 | 40  | Accent-aware STT, cultural context, translation | translation_support.py                          |
| 19: Ambient notifications    | 134-143 | 40  | Smart nudges, context timing, DND               | ambient_notifications.py                        |
| 20: Advanced analytics       | 144-153 | 40  | Trends, team performance, org learning          | advanced_analytics.py                           |

Totals: 20 phases, 153 days, ~612 hours, \$480K-720K (4 engineers × 6 months).

## Chapter 4 — The Coding Execution Policy

This is the governed coding policy. Every phase MUST follow it. The auditor verifies compliance by execution (P31). The policy has 3 stages: pre-phase (blocking), during-phase (per commit), and post-phase (blocking).

### 4.1 Pre-Phase (BLOCKING — before any code)

| Step | Action                                             | Principle                     |
|------|----------------------------------------------------|-------------------------------|
| 1    | Paste complete 8-field read receipt                | Gate (CEO rejects without it) |
| 2    | Verify prior phase gate passes by execution        | P1, P31                       |
| 3    | List files + grep call sites for parameter changes | P20                           |
| 4    | Audit every capture path for consent gating        | Ethical line                  |

### 4.2 During-Phase (per commit)

| Step | Action                                                             | Principle    |
|------|--------------------------------------------------------------------|--------------|
| 1    | Commit includes VERIFICATION section with pasted output            | P23          |
| 2    | grep call sites; M of N = (M/N)% done                              | P20, Gate 15 |
| 3    | Write 2 tests: unit + integration (production path)                | P22          |
| 4    | Re-run SSO scenario after shared-component changes                 | P24, P29     |
| 5    | Gate every confidence value on denominator (< 10 = "insufficient") | P25          |
| 6    | Execute consent-denied path; confirm zero capture                  | Ethical line |

### 4.3 Post-Phase (BLOCKING — before next phase)

| Step | Action                                                     | Principle |
|------|------------------------------------------------------------|-----------|
| 1    | Run phase gate commands; paste output                      | P1        |
| 2    | Run L0 gate; verify no regression                          | P29, P30  |
| 3    | Push to origin/main; paste HEAD + origin/main (must match) | Gate 11   |
| 4    | Auditor independently verifies by execution                | P31       |
| 5    | Only "Phase N PASS" allows next phase                      | The loop  |

### 4.4 Forbidden Anti-Patterns (auditor rejects on sight)

| Anti-pattern                                 | Principle violated |
|----------------------------------------------|--------------------|
| Theater tests (assert True, weak isinstance) | P27                |

| Anti-pattern                                   | Principle violated |
|------------------------------------------------|--------------------|
| Self-certification (✓ VERIFIED without output) | P5, P23            |
| Wiring gaps (0 of N callers pass new param)    | P20, Gate 15       |
| Stale-clone auditing (HEAD != origin/main)     | Gate 11            |
| Single-input testing (only the golden case)    | P28                |
| Confidence without denominator                 | P25                |
| Stealth framing (hide from meeting software)   | Ethical line       |

## Chapter 5 — Ambient Phase Details (9-20)

Each ambient phase has deliverables, an ethical guard, and a gate. The gates are verified by execution (P31). The ethical guards extend the Live Copilot bright line to the always-on layer: consent for every capture, no manipulation, no surveillance, no inference without confirmation.

| Phase | Deliverable                                | Ethical Guard                                                                          | Gate                                                          |
|-------|--------------------------------------------|----------------------------------------------------------------------------------------|---------------------------------------------------------------|
| 9     | Calendar awareness + commitment escalation | Calendar data only; no content with test cases; explicit consent only                  | test_calendar_explicit_test_only; commitment_escalation.py    |
| 10    | Sentiment & emotion tracking               | For user awareness only; never shown to other party                                    | test_sentiment_engine.py (emotion never shown to other party) |
| 11    | Deal health score                          | Score with denominator (P25); < 10 deals insufficiently (confidence has denominator)   | test_deal_health.py (confidence has denominator)              |
| 12    | Negotiation strategy                       | Preparation only; no manipulative tactics; evidence chain required                     | test_negotiated_strategy.py (evidence chain required)         |
| 13    | Relationship dynamics                      | Derived from org signals; no personal probes                                           | test_relationship_dynamics.py                                 |
| 14    | Cross-meeting threads                      | Threads by entity + topic; decisions tracked                                           | test_cross_meeting_threads.py                                 |
| 15    | Talk ratio + comms coach                   | Capability-building, not dominance-building                                            | test_talk_ratio_coach.py (no dominance suggestions)           |
| 16    | Meeting grade + analytics                  | Grade with denominator; follow-up tracked                                              | test_meeting_grade.py                                         |
| 17    | Email/Slack signal fusion                  | Per-channel consent, revocable; consent-denied path = 0                                | test_email_slack_signal_fusion.py (consent-denied path = 0)   |
| 18    | Multi-language support                     | Translation for user understanding; auto-translate support type (auto-translate gated) | test_multilingual_support.py (auto-translate gated)           |
| 19    | Ambient notifications                      | Respects DND + focus + off-hours; max 5/day                                            | test_ambient_notifications.py                                 |
| 20    | Advanced analytics                         | Team-level aggregate; no individual surveillance                                       | test_advanced_analytics.py (no individual PII)                |



## Chapter 6 — The Loop + Honest Disclosure

The governance loop is the enforcement mechanism. It is the reason the 9 code-quality findings from Audit 1 are actually fixed (verified by the auditor at HEAD 09b2b87), not just claimed fixed. The loop applies equally to all 20 phases:

| Step                 | Action                                                               | Who     |
|----------------------|----------------------------------------------------------------------|---------|
| 1. Before each phase | Read governance modules from disk; paste 8-field receipt             | Coder   |
| 2. During each phase | Gate P-numbers (P20, P22, P23, P24, P25) per commit                  | Coder   |
| 3. After each phase  | Run gate + L0; push; paste HEAD + origin/main                        | Coder   |
| 4. Auditor verifies  | Fetch → checkout → run gate independently (P31) → SSO scenario (P20) | Auditor |
| 5. Next phase        | Only after "Phase N PASS"                                            | Both    |

### 6.1 Honest Disclosure

This is a 153-day, 612-hour, \$480K-720K build. The coder can execute every CODE phase (1-20). The CEO must resource the GTM milestones (design partners, ARR, SOC2/HIPAA/GDPR) in parallel. The 20 phases deliver the product the CEO promised investors: an ambient organizational intelligence layer that works 24/7, learns from every interaction, and compounds institutional memory over time. This is not a feature. This is a category.

### 6.2 Read Receipt

GOVERNANCE LOOP READ RECEIPT (Coder):

- GOVERNANCE\_LOOP.md read at 2026-07-07T05:50:00Z. Key line: "THE LOOP CANNOT BE BROKEN."
- ENTROPY\_RECOVERY.md Part Four (P20-P26) read. P20: "If M of N call sites pass it, the fix is (M/N)% done." P26: "Principles don't enforce themselves. Re-application does."
- ENTROPY\_RECOVERY.md Part Five (P27-P34) read. P27: "Read at least the key assertions." P34: "The auditor's method is not a memory — it's a checklist."
- AUDITOR\_GOVERNANCE.md Gates 15-20 read. Gate 15: "A function signature with a new parameter proves nothing." Gate 17: "Reading the dedup logic proves nothing; executing it does."
- audit\_scripts/audit\_gates.sh read — enforces Gate 11 + full suite + all verify scripts.

Auditor: Read GOVERNANCE\_LOOP.md, AUDITOR\_GOVERNANCE.md, and ENTROPY\_RECOVERY.md Part Four + Part Five from disk before auditing. Do not audit from memory. The loop cannot be broken.