

Real-Time Meeting Intelligence, Done Right

An 8-phase, 33-day plan for the Maestro Live Copilot browser extension.
Side panel, not overlay. Consent-first. Evidence-grounded.
The anti-Cluely.

Document	Maestro Live Copilot — Governed Execution Roadmap
Audience	CEO, Engineering, Auditor
Date	2026-07-07
Current HEAD	e867cb7 (origin/main)
L0 gate	PASS (3/4 verified; L0.4 deferred)
Source	3-scene spec (Pre-Call / Live / Post-Call)
Ethical line	Side panel, consent-first, audit-logged — NOT Cluely

GOVERNANCE LOOP READ RECEIPT (Coder) — 2026-07-07T05:35:00Z
GOVERNANCE_LOOP.md: "THE LOOP CANNOT BE BROKEN." · P20: "If M of N call sites pass it, the fix is (M/N)% done." · P26: "Principles don't enforce themselves. Re-application does." · P27: "Read at least the key assertions of the tests you're counting." · P34: "The auditor's method is not a memory — it's a checklist."
> Auditor: Read GOVERNANCE_LOOP.md, AUDITOR_GOVERNANCE.md, and ENTROPY_RECOVERY.md Part Four + Part Five from disk before auditing.

Table of Contents

Table of Contents	2
Chapter 1 — The Non-Negotiable Ethical Line	3
Chapter 2 — The 3 Scenes (Target Experience)	4
2.1 Scene 1: Pre-Call Intelligence	4
2.2 Scene 2: Live Intelligence	4
2.3 Scene 3: Meeting Intelligence Captured	4
Chapter 3 — L0 Prerequisite Gate	5
Chapter 4 — The 8-Phase Plan	6
Chapter 5 — Technical Spec	7
Chapter 6 — Commit Format (P23)	7
Chapter 7 — The Governance Loop	8
7.1 Read Receipt	8
7.2 Honest Disclosure	8

Chapter 1 — The Non-Negotiable Ethical Line

Maestro Live Copilot is the anti-Cluely. The differentiator is not a feature — it is the ethical line. The following code paths are FORBIDDEN. If any appear, the auditor rejects the entire phase and the CEO is notified.

Forbidden path	Why it is rejected
Hiding from meeting software as a selling point	Stealth framing violates the bright line
Bypassing screen-share detection	Deception; enterprise buyers will not tolerate it
Unconsented recording	Consent MUST precede any audio capture (ConsentManager.checkConsent)
Exam/interview assistance modes	Manipulation; rejected absolutely by the constitution
"Undetectable mode" or stealth framing	The product is transparent — side panel, not overlay
Automatic answer injection into calls	Suggestions display in the panel; the user speaks them

What Maestro Live Copilot IS: side panel (not overlay), consent-first (every capture gated), audit-logged (every suggestion/capture/dismiss recorded), evidence-grounded (every suggestion cites organizational data). This is how enterprise products work.

Chapter 2 — The 3 Scenes (Target Experience)

2.1 Scene 1: Pre-Call Intelligence

Trigger: user opens a Google Meet / Zoom / Teams lobby URL; the Maestro extension detects the meeting metadata and surfaces a pre-call briefing in the side panel.

Surfaces: (1) Meeting context card — ARR at risk, renewal countdown, account health (from SituationSnapshot, which now has 27 fields per L0-1). (2) Attendee intelligence per person — interaction count from OEM signal history, commitment status from CommitmentTracker, last-interaction gap. (3) Suggested talking points — each citing the organizational data behind it ("12 interactions" links to signal history; "Day 45 of 60" links to the commitment).

The killer detail: every suggestion cites organizational data, not generic LinkedIn bios. This is the difference between Cluely and Maestro.

2.2 Scene 2: Live Intelligence

Trigger: user clicks "Start Copilot" in the side panel AFTER consent; audio capture begins via the offscreen document.

Four card types appear in real time, each color-coded:

Card	Color	Trigger	Content
Objection detected	Rose #FF5577	Transcript matches objection responses	only those cite validated runtimes; confidence bar from pattern
Commitment detected	Amber #FFB84D	Transcript matches commitment	Dedup against CommitmentTracker (Day X/Y, not a duplicate)
Organizational whisper	Purple #7C5CFF	Entity matches a GitHub PR / Slack channel	Stock-y confidence chain; confidence from source count
Historical pattern match	Cyan #5CC8FF	Conversation resembles a past meeting	Context cited; confidence from match count

The new-card animation is cardSlideIn (400ms ease-out) with a glow effect (box-shadow 0 0 20px rgba(255,214,10,0.15)) that fades after 5s. The user's eye is drawn to new suggestions immediately. Live transcript shows the last 3 chunks with speaker labels and trigger words highlighted.

2.3 Scene 3: Meeting Intelligence Captured

Trigger: call ends (WebSocket disconnect or tab close). The side panel shows a green checkmark and "Summary ready."

Surfaces: (1) Hero summary card — title, duration, participant count, transcript chunk count. (2) Key stats grid — commitments, objections, suggestions. (3) Commitments tracked — each with actor, Day X/Y, dedup status. (4) Objections raised — with response pattern. (5) Draft follow-up email — pre-written, citing specific commitments + patterns; copy / edit / open-in-Gmail. (6) What Maestro learned — new signals ingested, pattern data-point count, law-promotion threshold.

The killer detail: the "What Maestro learned" section shows the user the system gets smarter with every call. "The pricing objection pattern now has 4 data points — one more and it becomes a validated organizational law." This is the moat: judgment lock-in.

Chapter 3 — L0 Prerequisite Gate

The L0 gate MUST pass before Phase 1 begins. The L0 prerequisites ensure the shared OEM substrate is ready: SituationSnapshot has all 27 fields, OutcomeLedger is durable and tenant-scoped, the classifier handles natural language, and the SSO scenario still passes.

Gate	Verification	Status at e867cb7
L0.1 SituationSnapshot 27 fields	fields=27 missing=0/17	PASS
L0.2 OutcomeLedger functional	methods=[append,clear,close,count,get_all]	PASS
L0.3 Classifier new types	tentative/sarcasm/artifact (3/3)	PASS
L0.4 SSO scenario 7/7	pytest test_hybrid_e2e.py test_sprint_completion	DEEP RED

If any L0 check fails, STOP. Fix the L0 prerequisite first. The L0 gate is verified by execution (P1, P31), not by reading code or trusting commit messages.

Chapter 4 — The 8-Phase Plan

The plan is 33 days, ~132 hours. Each phase has deliverables, a gate, and a verification command. The gate MUST pass before the next phase begins. The auditor independently verifies each gate by execution (P31).

Phase	Days	Hours	Deliverable	Gate
1: Extension scaffold	1-3	12	manifest.json, background.js, consent-extension.js, onScreen.js, offscreen.js	Extension loads on chrome://extensions/
2: Audio + transcription	4-7	16	offscreen.js audio capture, Whisper transcription, speaker diarization	Transcription appears in chrome://extensions/
3: Scene 1 pre-call	8-11	16	lobby detection, pre_call.py, attendee identification, suggested topics	Pre-call screen suggests topics
4: Scene 2 live	12-18	28	live_engine.py, 4 card types, color borders, confidence bars, annotations, debug	Object confidence bars on annotations, debug
5: Scene 3 post-call	19-23	20	post_call.py, summary, stats, commitments, drafts, email, WhatsApp suggested notes	Outcome Ledger; drafts
6: Evidence + confidence	24-27	16	evidence-chain links, confidence display gate (P25), live evidence visibility	Confidence without evidence
7: Accessibility + polish	28-30	12	keyboard nav, aria-live, contrast >= 4.5, reduced-motion, 90 FPS, File Explorer	Lighthouse a11y, 90 FPS, File Explorer
8: Integration + audit	31-33	12	E2E test, cross-surface coherence (P24), E2E deployed on auditor test	E2E passes on auditor test; auditor says "Phase 8 PA"

Total: 33 days, ~132 hours. The plan is sequential — each phase depends on the prior. The governance loop enforces this: no phase begins until the prior phase's gate passes independent verification.

Chapter 5 — Technical Spec

The technical spec is taken directly from the source 3-scene description. These values are not suggestions — they are the design contract.

Element	Value
Side panel width	380px (Chrome native side panel)
New-card animation	cardSlideIn 400ms ease-out
Glow effect	box-shadow 0 0 20px rgba(255,214,10,0.15), fades after 5s
Rose (objection)	#FF5577
Amber (commitment)	#FFB84D
Purple (whisper)	#7C5CFF
Cyan (pattern)	#5CC8FF
Green (tracked)	#00D4AA
Typography	Inter (sans) + JetBrains Mono (numbers)
Accessibility	keyboard-navigable, focus-visible, aria-live regions
Consent	ConsentManager.checkConsent() gates every getUserMedia/getDisplayMedia
Audit log	Every suggestion, capture, dismiss logged

Chapter 6 — Commit Format (P23)

Every commit MUST include a VERIFICATION section with pasted command output. Claims without output are not evidence (P23). If the commit has no VERIFICATION section, the auditor rejects it.

Format:
feat(copilot): Phase N — short description

VERIFICATION:
\$ command
output

\$ command
output

Governance: P1 (execute), P23 (commit cites output), P26 (re-read from disk).
Read receipt: pasted in worklog at timestamp.

Chapter 7 — The Governance Loop

The governance loop is the enforcement mechanism. It is the reason the 9 code-quality findings from Audit 1 are actually fixed (verified by the auditor at HEAD 09b2b87), not just claimed fixed. The loop applies equally to the Live Copilot build:

Step	Action	Who
1. Before each phase	Read GOVERNANCE_LOOP.md + ENTROPY_RECOVERY Part Four from disk; paste 8-field read receipt	Coder
2. During each phase	Give the P-number principle each fix satisfies (P20 callers, P22 production path, P23 commit output, P24	Coder
3. After each phase	Run the phase gate commands; paste output in the commit; push to origin/main	Coder
4. Auditor verifies	Fetch → checkout HEAD → run gate independently (P31) → run Scenario (P29) → publish verdict	Auditor
5. Next phase	Only after auditor says "Phase N PASS" may the next phase begin	Both

7.1 Read Receipt

GOVERNANCE LOOP READ RECEIPT (Coder):

- GOVERNANCE_LOOP.md read at 2026-07-07T05:35:00Z. Key line: "THE LOOP CANNOT BE BROKEN."
- ENTROPY_RECOVERY.md Part Four read at 2026-07-07T05:35:00Z.
P20: "If M of N call sites pass it, the fix is (M/N)% done."
P26: "Principles don't enforce themselves. Re-application does."
- ENTROPY_RECOVERY.md Part Five read at 2026-07-07T05:35:00Z.
P27: "Read at least the key assertions of the tests you're counting."
P34: "The auditor's method is not a memory — it's a checklist."
- AUDITOR_GOVERNANCE.md Gates 15-20 read at 2026-07-07T05:35:00Z.
Gate 15: "A function signature with a new parameter proves nothing."
Gate 17: "Reading the dedup logic proves nothing; executing it does."
- audit_scripts/audit_gates.sh read — enforces Gate 11 + full suite + all verify scripts.

Auditor: Read GOVERNANCE_LOOP.md, AUDITOR_GOVERNANCE.md, and ENTROPY_RECOVERY.md Part Four + Part Five from disk before auditing. Do not audit from memory. The loop cannot be broken.

7.2 Honest Disclosure

This roadmap is the governed source of truth for the Maestro Live Copilot build. It covers all 3 scenes (pre-call, live, post-call), the 8-phase execution plan, the L0 prerequisite gate, the technical spec, and the governance loop. The coder can execute every phase. The auditor independently verifies each gate by execution (P31). The CEO enforces the loop: no message without a read receipt, no claim without a pushed commit, no gate the auditor hasn't independently verified. The loop cannot be broken.