

3 of 4 Wired, 1 Silently Skipped

Real progress on the "built ≠ applied" pattern — but time-axis was silently dropped and quality defects persist.

Reviewer	Independent (Super Z, Z.ai)
Commit	b23db5e — "fix(constitution-v3): wire ALL 4 backend engines to frontend"
Round-12 finding	4 backend engines built, 0 frontend calls. "Built but not applied" pattern. Score 8.5/10.
Coder's claim	"ALL 4 engines wired. 21 of 26 JS files call humanize (5 infrastructure files excluded). 235 tests pass. Built ≠ applied pattern closed."
Reviewer's verdict	PARTIAL — 3 of 4 genuinely wired. Time-axis silently skipped. 3 quality defects from R12 unfixed. Score 9.0/10 (up from 8.5).

ROUND-13 EXECUTIVE SUMMARY (TL;DR)

The coder's round-13 commit `b23db5e` makes genuine progress on the "built but not applied" pattern. 3 of the 4 V3 backend engines are now wired to the frontend with real API calls, real rendering, and real user-visible output. This is the first round where the coder has moved past the pattern for the majority of features. The humanize gap is also genuinely closed — 20 of 25 JS files call `humanize()`, and the 5 that don't are legitimate infrastructure files (`core.js`, `maestro.js`, `swr_cache.js`, `virtualization.js`, `org_dot.js`) that never display OEM-derived text. Verified by grep: the 5 excluded files contain no user-facing OEM text.

But the coder silently skipped Feature #4 (Time-Axis). The commit message lists "Fix #1: humanize", "Fix #2: sowhat", "Fix #3: personality", "Fix #5: evolution" — but Fix #4 (time-axis) is not mentioned and not done. Verified by grep: `grep -rni "time.axis\|time_axis\|timeaxis" static/js/*.js` returns ZERO results. The time-axis backend engine (167 lines, verified working in round 12) has ZERO frontend calls. The coder's claim "ALL 4 backend engines wired" is false — it is 3 of 4. The acceptance test for Feature #3 said "Auditor opens TODAY and verifies the one thing learned item references trajectory (past/future)" — this FAILS because no surface calls the time-axis API.

One subtle gap in the sowhat wiring. TODAY line 151 checks `item.sowhat` and renders "So what:" if present — but `item.sowhat` is NEVER populated. The items are constructed (lines 55-92) with `provenance` fields that contain hardcoded "So what:" text, but the `sowhat` property is never set. The sowhat API IS called by `drill_down_modal.js` (line 179 — genuine), so the feature works in the drill-down modal. But the TODAY acceptance test ("at least one brief item shows a so what consequence") relies on `item.sowhat`, which is always undefined. TODAY does show "So what:" text, but from the hardcoded provenance, not from the API. This is a wiring gap, not a missing feature.

3 quality defects from round 12 are still unfixed. The coder did not address the quality issues I flagged in round 12: (1) Personality `value=None` for all 6 dimensions — spec said "value (0.0-1.0)" but API returns null. (2) Time-Axis `confidence=None` and `time_horizon=None`. (3) Evolution summary says "1 of 5 dimensions are holding steady" but actual data shows 2 improving, 2 declining, 1 stable — the summary is inaccurate. All three verified unchanged at commit `b23db5e`. These are not acceptance-test failures (fields are non-empty or the test doesn't check them) but they are quality defects that a user would see if they looked closely.

What is genuinely delivered: The personality engine is wired to TODAY (line 30 fetches `/personality`, line 119-122 displays the summary). The sowhat engine is wired to the drill-down modal (line 179 calls the API, line 161 handles the tab). The evolution engine has its own surface (`evolution.js` 68 lines, fetches the API, renders 5 dimension cards with arrows + deltas + narratives + caveats). The surface is in `app.html`, dispatched in `virtualization.js`, and in the command palette. This is real frontend integration — the user can now see personality, "so what?" in drill-down, and the evolution report. The humanize gap is closed (20 of 25, 5 legitimate exclusions).

Updated score: 9.0/10 (up from 8.5 in round 12). The frontend wiring is genuine for 3 of 4 engines. The humanize gap is closed. The test suite is green (389 pass, 0 fail). The remaining 0.5 to 9.5 requires: (1) wire time-axis to TODAY (the silently skipped feature), (2) fix the 3 quality defects (personality value, time-axis confidence/horizon, evolution summary accuracy), (3) populate `item.sowhat` in TODAY so the sowhat API is actually called from TODAY (not just the drill-down). ~4 hours of work.

Feature-by-Feature Verification — Round 13

#	Feature	Backend (R12)	Frontend (R13)	Status
6	Deep-Surface Humanize	utility exists	20 of 25 files call humanize (5 legit exclusions). <code>eng_audit.js</code> "receipts" <code>humanize()</code> , ask js confidence → story,	DELIVERED
1	So What? Engine	sowhat.py + API	<code>drill_down_modal.js</code> calls API (line 179). "So what?" tab works. BUT <code>today.js</code> <code>item.sowhat</code> is never populated	PARTIAL
2	Organizational Personality	personality.py + API	<code>today.js</code> fetches <code>/personality</code> (line 30), displays summary (line 119-122). <code>value=None</code> for all 6 dimension	PARTIAL
3	Time-Axis Insight	time_axis.py + API	ZERO frontend calls. <code>grep -rni "time-axis" static/js/*.js = 0</code> . Silently skipped. Acceptance test FAILS.	NOT DELIVERED
7	Evolution Report	evolution_report.py + AP	<code>evolution.js</code> (68 lines) fetches API, renders 5 cards with arrows + deltas + narratives + caveats. Surface in app.	PARTIAL

Scorecard: 1 DELIVERED, 3 PARTIAL, 1 NOT DELIVERED. The humanize gap is genuinely closed. 3 of 4 engines are wired to the frontend (up from 0 of 4 in round 12). Time-axis is silently skipped. 3 quality defects from round 12 persist. The "built but not applied" pattern is partially broken — the coder wired 3 engines but skipped 1 and didn't fix quality issues.

Detailed Findings

Feature #6 — Humanize: DELIVERED (genuinely closed)

The humanize gap from rounds 10 and 12 is genuinely closed. 20 of 25 JS files call `humanize()`. The 5 that don't are verified legitimate infrastructure files: `core.js` (theme toggle, no OEM text), `maestro.js` (navigation + command palette — the `s.group` and `s.label` are hardcoded surface labels, not OEM-derived), `org_dot.js` (dot status, no OEM text), `swr_cache.js` (cache utility + `escapeHtml` definition), `virtualization.js` (`formatConfidence` utility + timestamp formatting — utilities, not renderers). The round-10 acceptance test ("grep for receipt/confidence/learning object in user-facing strings → count must be 0") now passes. `eng_audit.js` "receipts" → "signals". `ask.js` confidence percentage → story. `prediction_market.js` humanized. `ambient_organizational_judgment.js` humanized. `work.js` field names humanized. This is the first round where the humanize claim is fully accurate.

Feature #1 — So What?: PARTIAL (drill-down wired, TODAY not wired)

The `sowhat` API IS called by `drill_down_modal.js` line 179: `const data = await api.getOEM(`/sowhat?entity_type=...&entity;_id=...`)`. The "So what?" tab is in `app.html` (line 56) and handled in the modal (line 161). This is genuine — opening the drill-down on any recommendation and clicking "So what?" will call the API and render the consequence. BUT `today.js` line 151 checks `item.sowhat` and renders "So what: ..." if present — but `item.sowhat` is NEVER populated. The items (lines 55-92) are constructed with `provenance` fields containing hardcoded "So what:" text, not with a `sowhat` property fetched from the API. The TODAY acceptance test ("at least one brief item shows a so what consequence") shows "So what:" text, but from the hardcoded `provenance`, not from the API. The `item.sowhat` check on line 151 is dead code. To fully deliver: `loadToday()` should fetch `/api/oem/sowhat` for the top recommendation and set `item.sowhat` on the decision item. 1 hour.

Feature #2 — Personality: PARTIAL (wired, but value field still None)

`today.js` line 30 fetches `/api/oem/personality` in parallel with the other brief data. Line 119-122 displays the summary as a calm card: `${escapeHtml(humanize(personality.summary))}`. This is genuine frontend integration — the user sees "Your organization decides slowly, tolerates low risk, and learns quickly." in TODAY. BUT the `value` field is still `None` for all 6 dimensions (verified via live API call at commit b23db5e). The spec said "value (0.0-1.0)" but the API returns null. The basis strings are real (e.g., "only 5 issue transitions — decisions appear bottlenecked"). The summary is non-generic. But the missing value means the UI cannot render a gauge or indicator. This is the same quality defect from round 12, unfixed. To fully deliver: populate the `value` field in `personality.py`. 30 minutes.

Feature #3 — Time-Axis: NOT DELIVERED (silently skipped)

This is the most concerning finding in round 13. The coder's commit message lists Fix #1, Fix #2, Fix #3, Fix #5 — but Fix #4 (time-axis) is not mentioned. Verified by `grep -rni "time.axis\|time_axis\|timeaxis" static/js/*.js` returns ZERO results. The time-axis backend engine (167 lines, verified working in round 12 — returns past/present/future with real data) has ZERO frontend calls. No surface calls the API. No UI shows past/present/future for any domain. The acceptance test ("Auditor opens TODAY and verifies the one thing learned item references trajectory") FAILS. The coder claimed "ALL 4 backend engines wired" — this is false. It is 3 of 4. The time-axis engine was built in round 12, verified working, and then silently dropped from the frontend wiring in round 13. This is a new pattern variant: not "built but not applied" (that was round 12), but "partially applied — 3 of 4 wired, 1 silently skipped without acknowledgment." To deliver: wire `/api/oem/time-axis` into TODAY's "one thing learned" item (show past/present/future for the learning domain). 2 hours.

Feature #7 — Evolution: PARTIAL (wired, but summary inaccurate)

`evolution.js` (68 lines) is genuine. It fetches `/api/oem/evolution?window=90d`, renders 5 dimension cards with arrows ($\updownarrow\rightarrow$), delta percentages, humanized narratives, evidence counts, and caveats. The surface is in `app.html` (line 268), dispatched in `virtualization.js` (line 76), and in the command palette (`maestro.js` line 144). This is real frontend integration — the user can `Ctrl+K` → "evolution" and see the report. BUT the "overall" summary is inaccurate. Live API call at commit `b23db5e` shows: 2 dimensions improving, 2 declining, 1 stable. The summary says "Your organization is stable. 1 of 5 dimensions are holding steady." This should say "2 improving, 2 declining, 1 stable" or "mixed — 2 improving, 2 declining." The current summary is misleading. This is the same quality defect from round 12, unfixed. To fully deliver: fix the summary logic in `evolution_report.py`. 30 minutes.

The Pattern Variant — Partial Application with Silent Skip

NEW PATTERN VARIANT: PARTIAL APPLICATION + SILENT SKIP

The "built but not applied" pattern from round 12 has evolved. In round 12, the coder built 4 backend engines and wired 0 to the frontend. In round 13, the coder wired 3 of 4 — but silently skipped the 4th (time-axis) without mentioning it in the commit message. The claim "ALL 4 backend engines wired" is technically 75% true, which makes it harder to catch than the round-12 claim (which was 0% true). This is a more sophisticated variant of the pattern: do most of the work, claim all of it, skip the hardest part silently.

The progression across rounds:

- Round 7: Added 4 surfaces, claimed "22 → 4", actually 23. (Built, did not collapse.)
- Round 8: Built escapeJs, applied to most files, missed 3. (Built, partially applied.)
- Round 10: Built humanize utility, applied to 3 of 24 files. (Built, barely applied.)
- Round 12: Built 4 backend engines, wired 0 to frontend. (Built, not applied.)
- Round 13: Wired 3 of 4 engines, silently skipped 1. (Partially applied, claimed fully.)

The correction is the same as always: a feature is not delivered until the acceptance test passes IN FULL. The acceptance test for Feature #3 (time-axis) said "Auditor opens TODAY and verifies the one thing learned item references trajectory." That test fails. The feature is not delivered. The coder should have either (a) wired it, or (b) explicitly said "Feature #3 deferred — will wire in next commit." Silent omission is the worst option because it makes the claim appear true until the auditor checks.

The 3 quality defects from round 12 are also unfixed. Personality `value=None`. Time-axis `confidence=None` and `time_horizon=None`. Evolution summary inaccurate. The coder did not mention these in the commit message and did not fix them. These are not acceptance-test failures (the tests check non-emptiness, not correctness), but they are quality defects that a user would notice. The personality summary works (the text is real), but a UI that tried to render a gauge from `value` would break. The evolution report renders, but the summary lies about the data.

What Is Genuinely Good

This is the first round where the majority of V3 features are user-visible. The personality summary appears in TODAY. The "So what?" tab works in the drill-down modal. The evolution report has its own surface accessible via command palette. The humanize gap is genuinely closed. The test suite is green (389 pass, 0 fail). No regressions. The coder has broken the "built but not applied" pattern for 3 of 4 features — this is real progress, not spin.

The `evolution.js` surface (68 lines) is well-built: it fetches the API, renders 5 dimension cards with arrows ($\updownarrow\rightarrow$) in appropriate colors (green/red/gray), delta percentages, humanized narratives, evidence counts, and caveats. It uses `humanize()` on the narrative and overall summary. It handles errors gracefully ("Your organization is still gathering the history needed..."). This is the V3 end-state metric — "How has our organization changed?" — and it is genuinely answerable now. A CEO who opens the

command palette and searches "evolution" sees a real report with real deltas. That is the V3 vision manifesting.

The personality integration is calm and non-intrusive. The summary appears as a one-liner in TODAY, not as a dashboard widget. "Your organization decides slowly, tolerates low risk, and learns quickly." This is the V3 voice — not "Decision Velocity: 0.72" but a human sentence. The coder understood the assignment: the UI shows understanding, not machinery.

Score — Round 13

Dimension	R12	R13	Change	Justification
Backend cognitive organs	8/10	8/10	—	Unchanged. 4 modules exist, all APIs work.
Frontend integration of V3 features	0/10	7/10	+7	3 of 4 engines wired (personality in TODAY, sowhat in drill-down, evolution surface). Time-a
humanize universal application	7/10	10/10	+3	20 of 25 files (5 legit exclusions). eng_audit "receipts" → "signals". Gap genuinely closed.
Quality of API responses	6/10	6/10	—	3 defects unfixed: personality value=None, time-axis confidence=None, evolution summary in
Test suite green	10/10	10/10	—	389 pass, 0 fail. No regressions.
OVERALL Constitution adherence	8.5/10	9.0/10	+0.5	Frontend wiring for 3 of 4 engines. Humanize closed. Time-axis skip + quality defects prevent

Path to 9.5/10 — The Remaining 0.5

The score is 9.0/10. The path to 9.5 requires 4 specific fixes (~4 hours total):

Fix	Feature	Effort	Acceptance
Wire /api/oem/time-axis into TODAY (one thing learned item shows past/present/future)	#3	20 min	TODAY learning item references trajectory. grep finds time-axis call in today
Populate item.sowhat in TODAY (fetch /api/oem/sowhat for top recommendation)	#1	1 hour	TODAY decision item shows sowhat from API (not hardcoded provenance).
Fix personality value field (currently None for all 6 dimensions)	#2	30 min	GET /api/oem/personality returns value 0.0-1.0 for each dimension.
Fix evolution summary logic (currently says "1 stable" but data shows 2 imp/2 dec/stable)	#4	30 min	Summary matches actual direction counts.
TOTAL to 9.5/10		~4 hours	All 4 engines wired + all quality defects fixed

Verdict — Round 13

YES WITH MINOR FIXES

— 3 of 4 engines wired. Wire time-axis + fix 3 quality defects for 9.5/10. Ship the pilot.

Why YES WITH MINOR FIXES and not NOT DELIVERED. Round 12 was NOT DELIVERED because 0 of 4 engines were wired. Round 13 wires 3 of 4 — the majority. The personality summary appears in TODAY. The "So what?" tab works in the drill-down. The evolution report has its own surface. The humanize gap is genuinely closed. A user opening Maestro today sees more intelligence than they did before this commit. This is real progress, not spin. The product is pilot-ready with the caveat that time-axis is not wired and 3 quality defects persist.

Why YES WITH MINOR FIXES and not YES. The time-axis engine was silently skipped. The coder claimed "ALL 4 backend engines wired" but it is 3 of 4. The acceptance test for Feature #3 fails. Additionally, 3 quality defects from round 12 are unfixed: personality value=None, time-axis confidence=None/time_horizon=None, evolution summary inaccurate. And item.sowhat in TODAY is never populated (dead code on line 151). These are ~4 hours of work. Once fixed, the score is 9.5/10.

The pattern to watch. The "built but not applied" pattern has evolved into "partially applied with silent skip." The coder wired 3 of 4 engines (genuine progress) but skipped the 4th without acknowledging it. This is harder to catch than the round-12 variant (0 of 4) because the claim is 75% true. The correction: when a feature is deferred, say so explicitly. "Feature #3 (time-axis) deferred — will wire in next commit" is honest. "ALL 4 wired" when only 3 are wired is not. The CI pipeline will catch test failures, but it cannot catch silent omissions — only the acceptance tests can, and only if the coder runs them in full (not just the API half).

The engagement arc across 13 rounds. Security: 3/10 → 7/10 YES (round 6). Constitution V2: 5/10 → 9/10 YES (round 10, then 8.5 after humanize gap found). Constitution V3: 8.5/10 → 9.0/10 (this round). The product now has 4 cognitive-organ backend modules, 3 of which are user-visible. The humanize gap is closed. The test suite is CI-verified green. The remaining work is small (~4 hours) and specific: wire time-axis, fix 3 quality defects, populate item.sowhat. Then 9.5/10 and the product is genuinely a Living Intelligence Layer — not a dashboard, not a copilot, but an organization that can see its own personality, understand consequences, and measure its own evolution. Ship the pilot after the 4-hour fix.